

Travis Lilley  
Thinning Simulation

At the end of the second Exercise set, James claimed that thinning a Monte Carlo chain was unnecessary and in fact increased the variance of estimates. He asked us to run some simulations to test this claim ourselves. I ran several simulations using the Gibbs sampler from Exercise 2. My results confirm what James said: the posterior means for the  $\beta_i$  seem to have higher variance when we thin. The more dramatically we thin, the more this variance increases. R code is included at the end.

Simulation 1

Number of chains: 200

Chain length: 10000

Thin Ratio: 1/2

	$\hat{V}(\hat{\beta}_{0,\text{post}})$	$\hat{V}(\hat{\beta}_{1,\text{post}})$
Thin	$7.708 \cdot 10^{-9}$	$3.967 \cdot 10^{-6}$
No thin	$7.759 \cdot 10^{-9}$	$3.945 \cdot 10^{-6}$

Simulation 2

Number of chains: 200

Chain length: 10000

Thin Ratio: 1/5

	$\hat{V}(\hat{\beta}_{0,\text{post}})$	$\hat{V}(\hat{\beta}_{1,\text{post}})$
Thin	$6.947 \cdot 10^{-9}$	$4.056 \cdot 10^{-6}$
No thin	$6.566 \cdot 10^{-9}$	$3.558 \cdot 10^{-6}$

Simulation 3

Number of chains: 200

Chain length: 10000

Thin Ratio: 1/10

	$\hat{V}(\hat{\beta}_{0,\text{post}})$	$\hat{V}(\hat{\beta}_{1,\text{post}})$
Thin	$8.252 \cdot 10^{-9}$	$4.955 \cdot 10^{-6}$
No thin	$6.250 \cdot 10^{-9}$	$3.113 \cdot 10^{-6}$

Simulation 4

Number of chains: 500

Chain length: 10000

Thin Ratio: 1/2

	$\hat{V}(\hat{\beta}_{0,\text{post}})$	$\hat{V}(\hat{\beta}_{1,\text{post}})$
Thin	$6.489 \cdot 10^{-9}$	$3.585 \cdot 10^{-6}$
No thin	$6.397 \cdot 10^{-9}$	$3.548 \cdot 10^{-6}$

## R Code

```
# SDS 383D: Statistical Modeling II
# Exercise 2
# GDP Growth Gibbs Sampler
# Thinning simulation to see if thinning reduces variance

# Add libraries
library(mvtnorm)

# Load data
data <- read.csv("C:\\Users\\tr1l1\\Google Drive\\UT Austin\\SDS 383D Statistical
Modeling II\\Homework\\Exercises 2\\gdpgrowth.csv",
  header = T)

# Response
y <- data$GR6096

# Design matrix
X <- cbind(1, data$DEF60)

# Initialize prior parameters
d <- 1/100
eta <- 1/100
m <- c(0, 0)
K <- diag(c(0.001, 0.001))
h <- 1/100

# Number of Gibbs samples
G <- 10^5

# Number of Gibbs iterations
P <- 100

# Number of obs
n <- length(y)

# Objects to store Gibbs samples

# Lambda
L <- matrix(NA, nrow = G, n)

# Omega
w <- rep(0, G)

# Beta
B <- matrix(0, nrow = G, 2)

# Posterior mean of thinned betas
thin <- matrix(NA, nrow = 2, P)

# Posterior mean of non-thin betas
no_thin <- matrix(NA, nrow = 2, P)

# Initialize Gibbs
L[1, ] <- rep(1, n)
w[1] <- 1
B[1, ] <- c(0, 0)

# Run Gibbs M times
for(p in 1:P){
  for(i in 2:G){

    # Update location of sampler
    if(i %% (G/100) == 0){
      print(i)
    }

    # Update L
    for(j in 1:n){
      L[i, j] <- rgamma(1, h/2 + 1/2,
```

```

                                w[i-1]/2*(y[j]-x[j,] %**% B[i-1,])^2 + h/2)
}

# Diagonal matrix Lambda
L2 <- diag(L[i, ])

# Update w
w[i] <- rgamma(1, n/2 + d/2,
              (1/2)*(eta + t(y) %**% L2 %**% y + t(m) %**% K %**% m -
                t(t(X) %**% L2 %**% y + K %**% m) %**%
                solve(t(X) %**% L2 %**% X + K) %**%
                (t(X) %**% L2 %**% y + K %**% m))))

# Update B
B[i, ] <- rmvnorm(1, solve(t(X) %**% L2 %**% X + K) %**%
                  (t(X) %**% L2 %**% y + K %**% m),
                  solve(w[i] * t(X) %**% L2 %**% X + w[i] * K))
}

# Posterior mean of beta, no thin
no_thin[1, p] <- mean(B[, 1])
no_thin[2, p] <- mean(B[, 2])

# Thin
B_thin <- B[seq(1, G, by = 2), ]

# Posterior mean of beta, thin
thin[1, p] <- mean(B_thin[, 1])
thin[2, p] <- mean(B_thin[, 2])
}

```